Self-Correction Emerges from Recurrent Structure in a Minimal 35-Neuron Network: Ablation Evidence

Sungmoon Ong *softkorea@gmail.com*
Independent Researcher
Jinhae, South Korea
ORCID: 0000-0003-0215-6482

Abstract

Self-correction—the ability to revise one's own output through iterative refinement—is widely studied in both biological and artificial systems. While recent work has demonstrated that iterative refinement capabilities can develop in recurrent neural networks, the question of whether self-referential feedback is necessary for self-correction remains unaddressed. Here we construct a minimal recurrent network of 35 neurons and show that (1) self-correction develops when appropriate learning conditions are provided, even at minimal network scale; (2) ablating the recurrent self-referential loop eliminates self-correction entirely while preserving baseline pattern recognition; (3) replacing veridical self-feedback with distribution-matched shuffled feedback degrades performance below the no-feedback condition; (4) replacing self-feedback with the well-formed, in-distribution output of an independently trained clone model degrades performance substantially (gain = -0.059), ruling out marginal-distribution mismatch and demonstrating learned feedback-contract specificity—the recurrent weights are co-adapted to the specific model's output geometry; and (5) parameter-matched and compute-matched feedforward networks with equivalent or greater capacity show zero self-correction, arguing against simple capacity or depth explanations. Hyperparameter sweep across 80 configurations confirms that the phenomenon is broadly observed (68% of configurations under a lenient heuristic) within the explored grid. These findings suggest that self-referential structure appears to be a critical architectural factor for self-correction even at minimal scale, and that non-veridical feedback—whether noise or another model's valid output—is more harmful than no feedback at all. **Keywords**: self-correction, iterative refinement, output-feedback processing, feedback-contract specificity, ablation, recurrent neural network, minimal model, mechanistic interpretability

1 Introduction

1.1 The Inseparability Question

A central question in both cognitive science and artificial intelligence is whether iterative output-feedback processing—the capacity of a system to observe and revise its own output—can be dissociated from the system's general intelligence. In biological systems, reflective monitoring and task performance are deeply intertwined: damage to prefrontal monitoring circuits impairs not only reflective processing but also adaptive behavior more broadly (Fleming & Dolan, 2012). Whether artificial neural networks exhibit an analogous coupling is largely unexplored at the mechanistic level.

This question carries practical weight. If self-referential processing and task performance share neural substrates, constraining a model's capacity for output-feedback processing could carry hidden costs—a possibility that has not been experimentally tested.

1.2 Prior Work

Recent work has established that metacognitive capabilities can arise spontaneously in recurrent neural networks. Ma et al. (2025) demonstrated that uncertainty-monitoring behavior emerges without explicit training across 16 cognitive tasks in moment neural networks. Similarly, Li et al. (2025) showed that tiny recurrent neural networks of 1–4 units can discover cognitive strategies that model biological decision-making more accurately than classical cognitive models, with ablation studies revealing which architectural components are essential. In the domain of artificial life, evolutionary simulations have shown that metamemory based on self-reference to one's own memory emerges in neural networks with neuromodulation, producing structures that parallel established models of human metamemory (Yamato et al., 2022).

On the theoretical side, Kirsch & Schmidhuber (2022) have argued formally that self-referential architectures—where a network can modify or observe its own parameters—are representationally necessary for true meta-learning. This theoretical claim, however, has not been accompanied by empirical ablation evidence showing that removing the self-referential pathway eliminates self-correction.

In the domain of iterative computation, Graves (2016) introduced Adaptive Computation Time (ACT), and Banino et al. (2021) proposed PonderNet, both of which train recurrent networks to determine dynamically how many steps to compute before output. These systems achieve iterative refinement but require explicit halting mechanisms and specialized architectural modifications.

1.3 The Gap

Existing research has established that iterative refinement capabilities *can develop* in neural networks. What remains untested is the reverse question: *does removing the self-referential pathway degrade self-correction?* Furthermore, no study has isolated the contribution of model-specific output feedback from mere information flow, and no study has demonstrated self-correction in a minimal, fully transparent network where every neuron and connection can be directly inspected.

1.4 Contributions

This paper makes five contributions:

First, we demonstrate that self-correction develops in a 35-neuron recurrent network when two learning conditions are met: temporal freedom (no penalty for initial errors) and adequate gradient flow through the feedback pathway. This shows that the phenomenon is induced by learning environment rather than network scale, at least at this minimal scale.

Second, we show through selective ablation that removing the recurrent self-referential loop eliminates self-correction completely (correction gain drops to exactly zero) while preserving baseline pattern recognition (t=1 accuracy unchanged). This provides causal evidence that, in this architecture and task, self-correction depends on the recurrent self-feedback pathway.

Third, we introduce a novel diagnostic: the "fake mirror" test. By replacing veridical self-feedback with distribution-matched shuffled feedback, we demonstrate that corrupted self-reference is *more* harmful than absent self-reference (gain = -0.064 vs. 0.000). This suggests that the network is sensitive to the coordinate-structured identity of its own prior output, not merely to the statistical properties of an auxiliary signal.

Fourth, we extend the fake mirror test with a "clone feedback" condition: replacing self-feedback with the well-formed output of an identically trained model. This in-distribution feedback degrades performance substantially (gain = -0.059), ruling out marginal-distribution mismatch as the sole explanation and demonstrating learned feedback-contract specificity—the recurrent weights are co-adapted to the specific model's continuous output geometry, not merely to the statistical class of "valid feedback signals."

Fifth, we show through parameter-matched and compute-matched feedforward controls that the self-correction advantage cannot be explained by additional parameters or computational depth alone, and through a hyperparameter sweep across 80 configurations that the phenomenon is not obviously cherry-picked within the explored hyperparameter grid.

2 Methods

2.1 Architecture

We constructed a small recurrent multi-layer perceptron (RecurrentMLP) with four layers: input (10 neurons), hidden layer 1 (10 neurons), hidden layer 2 (10 neurons), and output (5 neurons), totaling 35 neurons. Hidden layers use ReLU activation; the output layer is linear. A recurrent connection feeds the previous output back to hidden layer 1 via a temperature-scaled hyperbolic tangent: feedback $= \tanh(y_{t-1}/\tau)$, where τ is the feedback temperature.

The network was deliberately kept small enough that all neurons and connections can be simultaneously visualized—a design decision motivated by the goal of mechanistic transparency. The entire system was implemented in pure NumPy without deep learning frameworks, ensuring that every computation is directly inspectable.

2.2 Task: Static Pattern Classification with Self-Correction

To isolate self-reference from memory, we used a static input task. Five prototype patterns were defined over 10 input dimensions: each class k activates dimensions $2k$ and $2k{+}1$ with amplitude 1.0, and adjacent class dimensions with amplitude 0.3 (inter-class ambiguity). Input samples were generated by adding Gaussian noise (σ = noise_level) to the prototype. On each trial, the network receives the same noise-corrupted pattern at every timestep $t = 1, 2, 3$. The task is 5-class classification. Because the input is identical at every timestep, the recurrent loop cannot serve as a memory pathway—its only function is to allow the network to observe and revise its own prior output.

This design cleanly separates two potential roles of recurrence: memory (carrying information about past inputs) and self-correction (carrying information about past outputs). In our task, only the latter is possible.

2.3 Training and the Conditions for Emergence

Training used full-batch gradient descent with learning rate 0.01 for 500 epochs on 200 training samples (200 held-out test samples), with $T = 3$ timesteps per trial. Backpropagation through time was implemented as a 3-step unroll, with numerical gradient checking (relative error threshold 10^{-4}) to verify correctness.

Two modifications proved necessary for emergence of self-correction:

Time-weighted loss. The per-timestep loss $L(t)$ is softmax cross-entropy between the output logits and the one-hot target. The total loss is weighted across timesteps as $L = w_1 \cdot L(t{=}1) + w_2 \cdot L(t{=}2) + 1.0 \cdot L(t{=}3)$. With uniform weighting ($w_1 = w_2 = 1/3$), the network optimizes for immediate accuracy at $t{=}1$ and has no incentive to develop an iterative correction strategy. Setting $w_1 = 0.0$ grants the network temporal freedom: "your initial guess carries no penalty; only the final answer matters." This creates optimization pressure to use the feedback loop for refinement.

Temperature scaling. Without temperature scaling, output logits quickly exceed ± 3, saturating the tanh feedback function and driving the gradient of the recurrent weights to near zero. Setting $\tau = 2.0$ prevents saturation and maintains gradient flow through the feedback pathway.

We report these modifications transparently because they illuminate a broader point: the failure of emergence under uniform loss and unscaled feedback is itself informative. It demonstrates that emergence depends on learning conditions, not latent capacity—the network had the architectural capacity for self-correction all along but could not discover it under conditions that penalized intermediate exploration.

2.4 Ablation Design

We employed nine evaluation conditions (Baseline plus eight ablation/control conditions) to isolate the contribution of self-referential feedback:

Baseline: The fully trained recurrent model, evaluated normally.

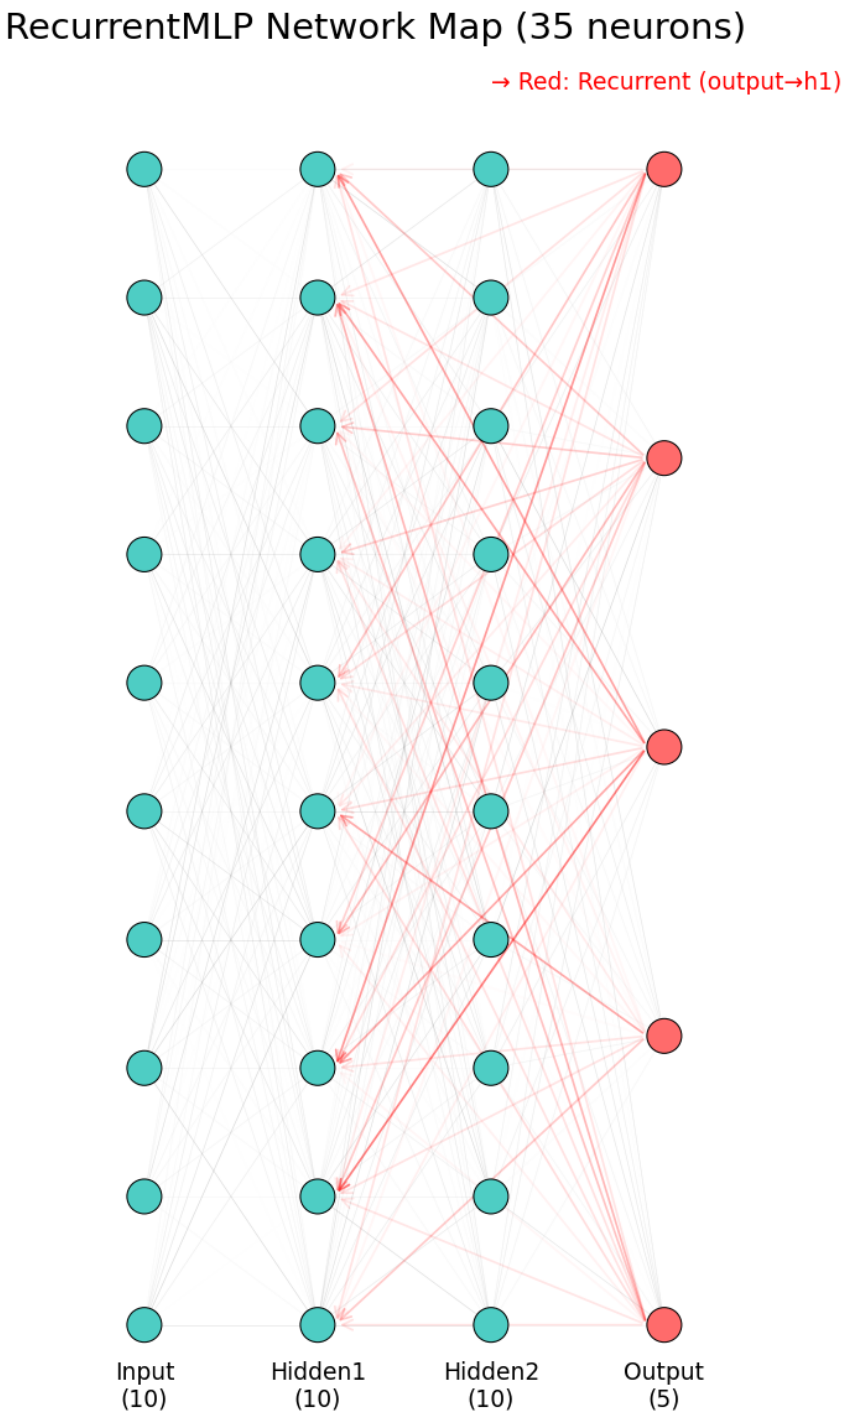

Figure 1: Architecture of the 35-neuron RecurrentMLP. The recurrent feedback pathway (output → hidden layer 1) is shown in red. All weights are directly accessible and inspectable.

Group A (Recurrent ablation): All recurrent weights set to zero post-training. The feedback pathway is structurally intact but carries no signal.

Group B1 (Random ablation): The same number of weights (50) randomly zeroed across all weight matrices. Repeated 30 times per model with different random seeds.

Group B2 (Structured ablation): The h2-to-output weight matrix zeroed entirely, removing a major feedforward pathway.

Group C1 (Permutation feedback): The recurrent connections are preserved, but the feedback vector is element-wise permuted at each timestep. This maintains the exact distribution, norm, and dimensionality of the feedback signal while destroying its coordinate-structured identity—the "fake mirror" condition.

Group C2 (Clone feedback): The recurrent connections are preserved, and the feedback is replaced with the output of a different model trained with identical architecture and procedure but a different random seed. Each target model (seed 0–9) is paired with a separately trained donor model (seed 100–109), ensuring full statistical independence between pairs. The clone's output is well-formed, in-distribution, and a valid classification output—but it is not the model's own. This condition controls for marginal-distribution mismatch: the clone's output has valid statistics (correct dimensionality, plausible magnitudes, trained output geometry), so any remaining degradation indicates trajectory-level specificity rather than distributional artifact. As a robustness check, we also tested a data-matched variant in which each donor model was trained on the exact same empirical data (identical noise realizations) as its target, differing only in weight initialization, to isolate the pure effect of initialization trajectory from any data-distribution confound.

Group D (Feedforward baseline): A network trained from scratch with no recurrent connections.

Group D′ (Parameter-matched feedforward): A feedforward network with an additional skip connection from input to output (50 weights), matching the total parameter count of the recurrent model. This addresses the possibility that the recurrent model's advantage is simply due to having more parameters.

Group D″ (Compute-matched feedforward): A deep feedforward network with 6 hidden layers of 10 neurons each (Input(10) $\rightarrow$ H1(10) $\rightarrow \ldots \rightarrow$ H6(10) $\rightarrow$ Output(5)), matching the computational depth of the recurrent model's 3-timestep unroll (3 timesteps $\times$ 2 hidden layers = 6 layer traversals). This has 715 parameters (vs. Baseline 325) and approximately 650 multiply-add operations (vs. Baseline $\sim$900 MACs across 3 timesteps). This addresses the possibility that the recurrent model's advantage comes from additional computation rather than from feedback content.

2.5 Neuron Importance Analysis

To quantify the contribution of individual neurons to pattern recognition and self-correction, we used decoupled single-neuron knockout analysis. For Hidden Layer 1 neurons (which receive both feedforward and recurrent input), intelligence and correction were measured separately: (a) intelligence importance was measured by zeroing only the feedforward incoming weights (W_{ih1} column + bias), observing the drop in t=1 accuracy (Δacc_{t1}); (b) correction importance was measured by zeroing only the recurrent incoming weights (W_{rec} column), observing the drop in correction gain (Δgain). This decoupling prevents a confound where feedforward-critical neurons appear artificially important for correction simply because destroying their feedforward pathway also collapses the correction baseline. For Hidden Layer 2 neurons (which have no direct recurrent input), full knockout was used for both measures; we note that the intelligence-correction confound may still exist for these neurons. This analysis was performed on a single trained model (seed=0, noise=0.5) to generate a per-neuron importance heatmap.

2.6 Statistical Design

All experiments were conducted across 10 independently initialized models (seeds 0–9). Groups B1 and C1 were repeated 30 times per model with different ablation/permutation seeds. Group C2 paired each target model (seed 0–9) with an independently trained donor model (seed 100–109). The primary measure was **correction gain**: accuracy at t=3 minus accuracy at t=1. Secondary measures included recurrent contribution norm ($\|\tanh(y_{t-1}/\tau) \cdot W_{\text{rec}}\|$) and step delta ($\|y_t - y_{t-1}\|$). For B1 and C1, the 30 repeats were averaged within each model before inference; the inferential unit is the model (N=10), not the repeat.

The pairing structure differs by group type: Groups A, B1, C1, and C2 are post-hoc interventions on the same trained model (within-model pairing); Groups D, D′, and D″ are separately trained seed-matched controls (matched-design pairing). B2 was excluded from the primary correction family a priori because its catastrophic feedforward degradation renders correction gain comparison uninformative. The secondary family (C1-vs-A and C2-vs-A) was pre-specified to test whether non-veridical feedback is worse than absent feedback.

Significance was assessed with exact Wilcoxon signed-rank tests on paired model-level means ($2^{10} = 1{,}024$ sign assignments enumerated), with Holm-Bonferroni correction for the primary family of comparisons. 95% bootstrap confidence intervals (10,000 resamples) are reported as descriptive summaries; the exact Wilcoxon p-values are the primary inferential quantities. Noise level was set to 0.5 for the primary analysis, with a full sweep across $[0.1, 0.2, 0.3, 0.5, 0.7, 1.0]$.

2.7 Robustness Analysis

To address potential concerns about parameter sensitivity, we conducted a hyperparameter sweep across 80 configurations: $w_1 \in \{0.0, 0.1, 0.2, 0.3\}$, $w_2 \in \{0.1, 0.2, 0.3, 0.5\}$, $\tau \in \{1.0, 1.5, 2.0, 3.0, 5.0\}$, with 10 independent models per configuration (800 total experiments). For the purpose of mapping the grid, we used a lenient heuristic to flag configurations showing emergence: mean gain > 0 with at least 60% of models showing positive gain. This criterion is deliberately liberal (a sign test at 6/10 yields $p \approx 0.38$) and is intended only for exploratory grid characterization, not as a formal significance test; the primary statistical claims rest on the exact Wilcoxon tests reported in Section 3.

3 Results

3.1 Emergence of Self-Correction

Under the reported hyperparameters ($w_1 = 0.0$, $w_2 = 0.2$, $\tau = 2.0$), the Baseline model achieved a correction gain of $+0.042 \pm 0.030$ (mean $\pm$ SD across 10 models), with a 95% bootstrap CI of $[+0.023, +0.059]$ that excludes zero. Accuracy improved from 0.698 ± 0.054 at t=1 to 0.740 ± 0.055 at t=3. This confirms that the 35-neuron network learned to use its own prior output to correct its predictions.

3.2 Ablation Results

Recurrent ablation (Group A): Correction gain dropped to exactly 0.000 ± 0.000. Accuracy at t=1 remained identical to Baseline (0.698 ± 0.054); this is architecturally guaranteed because the hidden state is initialized to zeros, making recurrent feedback exactly zero at t=1 regardless of ablation. This architectural isolation serves as a useful control: it confirms that the ablation targets only the correction mechanism, leaving the feedforward pathway intact. The exact Wilcoxon signed-rank test against Baseline yielded $p = 0.023$ after Holm-Bonferroni correction. The secondary measures corroborate this: Baseline recurrent contribution norm was 2.163 ± 0.311 and step delta was 1.162 ± 0.335, both dropping to exactly zero in Group A. We note that under static input, zero gain is structurally guaranteed when recurrence is removed, since no new information enters the network across timesteps. The non-trivial findings in this experiment are therefore (a) that Baseline gain is positive (the network *learned* to use feedback for correction) and (b) that C1/C2 gain is *negative* (non-veridical feedback causes active interference).

Random ablation (Group B1): Mean gain $= -0.004 \pm 0.016$. Accuracy at t=1 dropped substantially (0.515 ± 0.048), indicating general degradation. The correction mechanism was also lost, but this is confounded by the overall performance collapse.

Structured ablation (Group B2): The h2-to-output pathway ablation completely severed the hidden layers from the output, causing catastrophic degradation to chance-level accuracy (t=1 $= 0.199 \pm 0.032$). This serves as a sanity check rather than a nuanced structural ablation: it confirms that destroying the entire predictive pathway is qualitatively different from the targeted recurrent ablation in Group A. B2 was excluded from the Holm-Bonferroni correction family for this reason.

Table 1: Main results at noise level 0.5 (mean ± SD across N=10 models). Correction gain = acc(t=3) − acc(t=1). All p-values are Holm-Bonferroni corrected.

Group	acc(t=1)	acc(t=3)	Gain	95% CI	Corr. p
Baseline	$.698 \pm .054$	$.740 \pm .055$	$+\mathbf{.042}$	$[+.023, +.059]$	—
A (Recurrent cut)	$.698 \pm .054$	$.698 \pm .054$	.000	$[.000, .000]$	.023
B1 (Random cut)	$.515 \pm .048$	$.511 \pm .037$	$-.004$	$[-.013, +.006]$	.016
B2 (Structural cut)	$.199 \pm .032$	$.199 \pm .032$	.000	—	—
C1 (Shuffled fb.)	$.698 \pm .054$	$.634 \pm .041$	$-\mathbf{.064}$	$[-.095, -.036]$	.012
C2 (Clone fb.)	$.698 \pm .054$	$.639 \pm .081$	$-\mathbf{.059}$	$[-.101, -.015]$	.012
D (Feedforward)	$.746 \pm .067$	$.746 \pm .067$	.000	—	.023
D′ (Param-matched)	$.822 \pm .031$	$.822 \pm .031$	.000	—	.023
D″ (Compute-matched)	—	—	.000	—	—

Parameter-matched feedforward (Group D′): This model achieved the highest t=1 accuracy (0.822 ± 0.031), demonstrating that additional parameters do improve initial pattern recognition. However, correction gain was exactly 0.000 ± 0.000 (a structural certainty given the static input and lack of recurrence). Additional capacity without recurrent structure produces no self-correction whatsoever.

Compute-matched feedforward (Group D″): The 6-layer deep feedforward network, which matches the recurrent model's computational depth (6 layer traversals), achieved high single-pass accuracy but correction gain was exactly 0.000 ± 0.000 by construction (no recurrence, no feedback). This demonstrates that additional computational depth alone—without feedback content—cannot produce iterative correction. Combined with D′, this rules out both parameter count and computational cost as explanations for the recurrent model's self-correction advantage.

Table 1 summarizes all experimental conditions.

3.3 The Fake Mirror Effect

Group C1 (permutation feedback) produced a striking result. Despite preserving the full recurrent architecture, connection weights, and the statistical distribution of the feedback signal, correction gain was -0.064 ± 0.048—not merely zero, but significantly *negative*. The 95% CI was $[-0.095, -0.036]$. This means that shuffled feedback actively degrades performance below the initial prediction.

A natural objection is that permuted feedback constitutes an out-of-distribution (OOD) input that disrupts computation regardless of self-referential content. Group C2 (clone feedback) was designed to control for marginal-distribution mismatch. In C2, the feedback signal is the well-formed output of an independently trained model—same architecture, same data distribution and training procedure, different random seed. This feedback has valid marginal statistics (correct dimensionality, plausible magnitudes, trained output geometry). It is simply not the model's *own* output.

C2 produced substantial degradation: gain $= -0.059 \pm 0.069$, with a 95% CI of $[-0.101, -0.015]$ ($p = 0.012$ after Holm-Bonferroni correction; raw $p = 0.0020$). The degradation was comparable to C1, and the pattern held across intermediate noise levels (0.2–0.7). A direct paired Wilcoxon test between C1 and C2 was not significant ($p = 0.846$, $N = 10$), so we cannot claim C2 is statistically worse than C1—only that both are significantly worse than Baseline, and that in-distribution feedback is at least as damaging as shuffled feedback.

The gradient of degradation across feedback conditions is:

- **Group A** (no feedback): gain = 0.000. Neutral.
- **Group C1** (shuffled feedback): gain = −0.064. Harmful.
- **Group C2** (clone feedback): gain = −0.059. Comparably harmful.

To directly test the claim that wrong feedback is worse than no feedback, we performed paired Wilcoxon signed-rank tests comparing C1 and C2 against A, treated as a secondary family of two comparisons with Holm-Bonferroni correction. Both comparisons remained significant: C1 vs A (raw $p = 0.0039$, corrected $p = 0.0078$, 95% CI of difference $[-0.095, -0.036]$) and C2 vs A (raw $p = 0.0273$, corrected $p = 0.0273$, 95% CI $[-0.101, -0.015]$). Note that the C1 vs A CI is identical to C1's standalone CI because Group A's gain is strictly zero for all 10 models, so the paired difference reduces to C1's own distribution; the same logic applies to C2 vs A. This confirms that non-veridical feedback actively degrades performance below the no-feedback baseline.

That C2 degrades performance despite providing in-distribution feedback has a mechanistic explanation. The recurrent weights (W_{rec}) are co-adapted to the model's own specific feedforward feature detectors during training. Clone feedback, though well-formed, activates these co-adapted weights with a structured but misaligned signal, pushing the hidden state into a coherent but incorrect region of the model's latent space. Because the clone's output is a plausible classification vector—unlike shuffled feedback, which is partially incoherent—the recurrent pathway processes it as a reliable self-signal and follows it into systematically wrong corrections. We term this phenomenon *learned feedback-contract specificity*: the recurrent weights have learned a "contract" with the model's own output geometry that cannot be satisfied by another model's outputs, even when those outputs are statistically valid.

This result rules out marginal-distribution mismatch as the sole explanation for C1's degradation. C2's clone output has valid marginal statistics, yet produces comparable harm—indicating that the degradation operates at the level of trajectory-specific feedback alignment, not at the level of distributional validity.

We note as a caveat that the C1-vs-C2 comparison is not statistically significant ($p = 0.846$, $N = 10$), so we cannot distinguish the magnitude of C1 and C2 harm—only that both are significantly worse than Baseline and than the no-feedback condition. The absence of a C1-vs-C2 difference is an absence of evidence, not evidence of absence.

As a robustness check, we also tested a stricter variant in which the clone was trained on the exact same empirical data (identical noise realizations), isolating the pure effect of initialization trajectory. This data-matched clone still eliminated self-correction (gain $= -0.054$, $p = 0.006$ vs Baseline), though the comparison against the no-feedback condition (Group A) was marginally non-significant ($p = 0.080$), consistent with the expectation that shared training data induces partial representational overlap between models.

3.4 Noise Dependence

The correction gain varied with noise level, peaking at noise = 0.2 (gain $\approx +0.099$) and declining at both lower and higher noise levels. At low noise (0.1), feedforward processing alone was sufficient, rendering self-correction unnecessary. At high noise (0.7–1.0), the signal was too degraded for correction to recover. The optimal zone—where self-correction provides the greatest advantage—corresponds to intermediate difficulty, consistent with the cognitive science concept of a metacognitive "Goldilocks zone" where uncertainty is high enough to warrant reflection but low enough for reflection to be productive.

3.5 Neuron Overlap Analysis (Exploratory)

Decoupled neuron importance analysis of Hidden Layer 1 (performed on a single model, seed=0)—where feedforward and recurrent contributions can be cleanly separated—revealed three functional groups: neurons important primarily for initial pattern recognition (intelligence-specialized, e.g., h1_6), neurons whose recurrent input contributes primarily to self-correction (correction-specialized, e.g., h1_8), and neurons important for both (shared, e.g., h1_1, h1_5, h1_7). Notably, h1_9 showed *negative* correction importance (-0.035): removing its recurrent input *improved* correction gain, suggesting that not all recurrent channels contribute positively. For Hidden Layer 2 (where decoupling is not possible without timestep masking), full knockout was used; the intelligence-correction confound should be noted when interpreting these values.

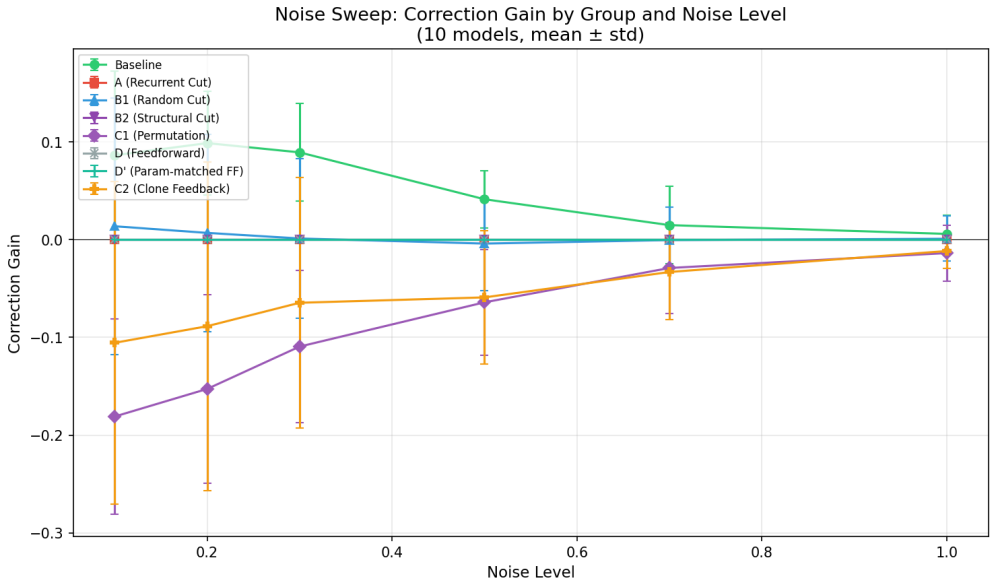

Figure 2: Correction gain across noise levels for Baseline and ablation groups. Self-correction peaks at moderate noise ($\sigma = 0.2$) and declines at both extremes.

3.6 Robustness

The hyperparameter sweep confirmed that self-correction emergence is broadly observed within the explored grid: 54 of 80 configurations (68%) met the lenient emergence heuristic defined in Section 2.7. The primary determinant was w_1 (t=1 loss weight): emergence occurred in 95% of configurations with $w_1 = 0.0$, 90% with $w_1 = 0.1$, 75% with $w_1 = 0.2$, and only 10% with $w_1 = 0.3$. Temperature τ was secondary: emergence was stable across $\tau \in [1.0, 3.0]$ and degraded only at $\tau = 5.0$. The t=2 weight w_2 had negligible effect across its full range.

The reported experimental configuration ($w_1 = 0.0$, $w_2 = 0.2$, $\tau = 2.0$) ranked 8th out of 80 configurations by mean gain (top 10%). The top configuration ($w_1 = 0.0$, $w_2 = 0.1$, $\tau = 1.0$) achieved gain = +0.054. This confirms that our results do not depend on cherry-picked hyperparameters.

4 Discussion

4.1 Emergence Depends on Conditions, Not Merely Scale

The prevailing narrative in AI research emphasizes scale as the driver of emergent capabilities (Wei et al., 2022a). Our results complicate this picture, at least in one minimal setting. Self-correction developed in a network of 35 neurons—orders of magnitude smaller than systems where iterative refinement is typically studied. In this toy model, what mattered was not the number of neurons but two specific learning conditions: temporal freedom (the absence of penalty for initial errors) and gradient flow through the feedback pathway. Whether this principle generalizes to larger architectures and more complex tasks remains an open question.

This finding resonates with the "grokking" phenomenon (Power et al., 2022), where sudden generalization in small networks was shown to depend on optimization dynamics rather than model size. Our result extends this observation to a different emergent property—self-correction through self-reference—though whether the parallel holds beyond this minimal setting remains to be tested.

A speculative implication—which we flag explicitly as such—is that self-correction-like behavior may be latent in recurrent architectures more broadly, requiring only appropriate training signals to manifest. The network in our experiment possessed the architectural capacity for self-correction from initialization—it

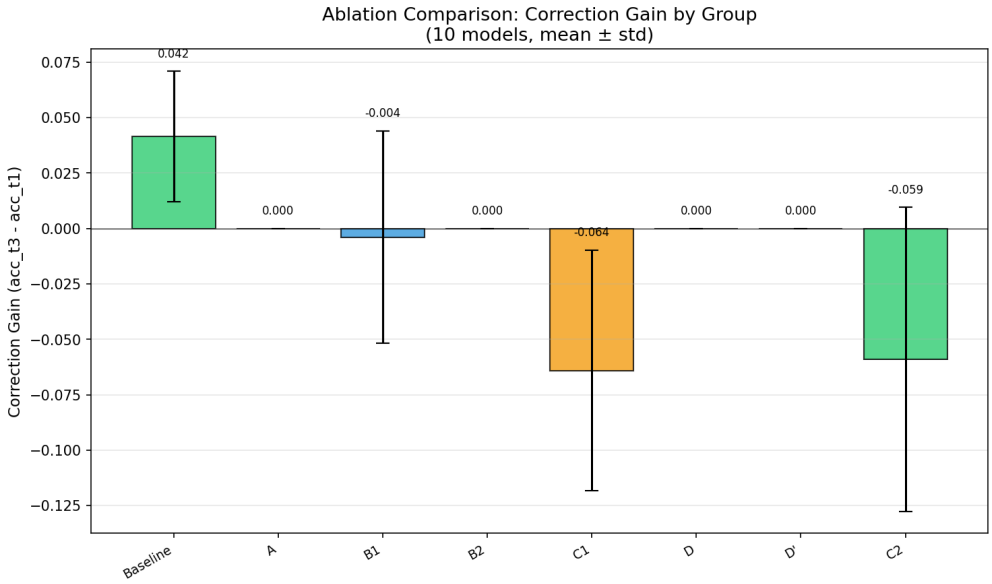

Figure 3: Correction gain by experimental group at noise level 0.5. Baseline shows positive gain; ablation groups show zero or negative gain. Error bars: ±1 SD across 10 models.

simply could not discover the strategy under conditions that punished intermediate exploration. Whether this observation generalizes beyond minimal static-input settings remains entirely open.

4.2 The Fake Mirror Is Worse Than No Mirror

The most novel finding is that non-veridical self-reference is more harmful than absent self-reference. This "fake mirror" effect manifests in two forms: shuffled feedback (C1, gain $= -0.064$) and clone feedback (C2, gain $= -0.059$), both significantly worse than no feedback (A, gain $= 0.000$). While the difference between C1 and C2 is not statistically significant (paired Wilcoxon $p = 0.846$), both are individually significant against Baseline (corrected $p = 0.012$), and the critical finding is that in-distribution feedback is at least as damaging as shuffled feedback.

The clone feedback result (C2) is particularly informative because it distinguishes two levels of distributional mismatch. C2 rules out *marginal-distribution* mismatch: the clone's output has valid statistics, correct dimensionality, and plausible magnitudes. What C2 reveals is *trajectory-level feedback-contract specificity*: the recurrent weights are co-adapted to the specific model's continuous output geometry, and substituting another model's equally valid output produces comparable harm. The only difference between Baseline and C2 is *whose* output the network reads back—its own, or that of an identically trained model with a different random seed. This difference alone reverses the sign of self-correction.

The mechanistic explanation is that each model's recurrent weights (W_{rec}) become co-adapted during training to the model's own specific continuous logit magnitudes and error signatures. Although the output space is class-aligned (both models map class k to index k, since the output layer is directly supervised), the continuous logit values—their magnitudes, relative ratios, and how they evolve across timesteps—differ between independently trained models due to different weight initializations. W_{rec} has learned to interpret these model-specific logit patterns to drive correction; clone feedback, though well-formed and in-distribution, provides logit magnitudes and ratios that W_{rec} was not trained to process, pushing the hidden state into an incorrect region. We term this *learned feedback-contract specificity*.

4.3 Implications for AI Systems

We emphasize at the outset that extrapolation from a 35-neuron static classification task to billion-parameter language models requires extreme caution. The architectural, task, and training differences are vast. With that caveat, the narrow factual finding is that in this minimal recurrent architecture, the network's ability to correct its own errors was entirely dependent on its ability to observe its own prior output. Remove the output-feedback pathway, and correction vanishes—even though the network retains its full feedforward capability. This is consistent with theoretical arguments that self-referential structure is representationally necessary for certain forms of meta-learning (Kirsch & Schmidhuber, 2022).

4.4 Connection to Biological Metacognition

The functional organization observed in our minimal network—with partially overlapping neurons serving both pattern recognition and self-correction—bears a surface-level resemblance to partially overlapping neural correlates of monitoring and control in prefrontal cortex (Morales et al., 2018). We note this structural parallel without claiming generality: the analysis was exploratory (single model, seed=0) and the toy setting precludes strong conclusions.

4.5 Relationship to Chain-of-Thought and Test-Time Compute

The time-weighted loss structure that enabled emergence in our experiment shares a conceptual link with recent advances in large language model reasoning. Chain-of-Thought prompting (Wei et al., 2022b) and test-time compute scaling both involve granting the model additional processing steps before requiring a final answer. Our time-weighted loss achieves an analogous effect at the training level: by removing penalty from early timesteps, we create pressure for the network to use those steps for exploratory computation rather than premature commitment.

Unlike PonderNet (Banino et al., 2021) and Adaptive Computation Time (Graves, 2016), our approach requires no explicit halting mechanism or specialized architecture—only a reweighting of the loss function. This simplicity may be of interest for understanding the minimal conditions under which iterative refinement emerges.

4.6 Limitations

Several limitations must be acknowledged.

First, our self-correction was *induced* by the time-weighted loss, which explicitly removes penalty for early errors. This is not spontaneous emergence—it is a learned behavior that requires specific training conditions. As noted in §3, Group A's zero gain is structurally guaranteed under static input, so the non-trivial claims rest on Baseline's positive gain (learned) and C1/C2's negative gain (active interference).

Second, our model is a toy system. Thirty-five neurons performing static classification cannot capture the complexity of language, planning, or open-ended reasoning. The value of our experiment lies in its transparency and causal clarity, not in its ecological validity.

Third, self-correction as operationally defined here (improving accuracy from t=1 to t=3 on the same static input) is a narrow proxy for the richer concept of self-awareness. We do not claim that our network "knows" it is correcting itself, only that it systematically uses its own prior output to improve.

Fourth, one might object that iterating a recurrent network for T=3 steps on static input is equivalent to passing through a weight-shared 3× deeper feedforward network, and that the performance improvement from t=1 to t=3 reflects increased computational depth rather than self-correction. Group A speaks to this: when recurrent weights are zeroed, the network produces identical output at every timestep (gain = 0.000), because with static input and no feedback, each step is a redundant re-computation of the same feedforward pass—no new information enters the network. This shows that the correction in the Baseline does not come from repeated computation per se, but from the novel information injected via the feedback pathway at each

step. Whether to call this "depth" or "self-correction" is partly terminological, but the key empirical fact is that the feedback content—not the number of processing steps—is necessary and sufficient for the gain.

Fifth, the parameter-matched feedforward model (Group D′) achieved higher absolute accuracy at t=1 (0.822) than the recurrent Baseline at t=3 (0.740). This indicates that for this specific static classification task, distributing parameters across a feedforward architecture is more efficient for single-pass pattern recognition. Our claim is not that recurrent self-correction achieves superior absolute performance, but that it represents a qualitatively different computational strategy—iterative refinement through self-reference—that is architecturally inseparable from the recurrent pathway. In tasks requiring genuine iterative reasoning under capacity constraints, this strategy may prove essential.

Sixth, the robustness analysis, while extensive (800 runs), explored only three hyperparameters. Other factors—network depth, activation function, learning rate schedule—may influence emergence in ways we have not characterized.

Seventh, our use of static inputs cleanly separates self-correction from memory, but in autoregressive language models and biological systems, the recurrent pathway must simultaneously manage temporal memory (past inputs) and self-reflection (past outputs). How these two functions compete for recurrent bandwidth—and whether one can be suppressed without degrading the other—remains an open question for future work.

Eighth, while Group D′ controls for parameter count and Group D″ controls for computational depth, neither control can produce correction gain by construction (no recurrence under static input). This means D′/D″ demonstrate that parameters and compute are *insufficient* for correction, but the argument that they are *unnecessary* (i.e., that feedback content is the sole driver) rests on the Group A ablation.

Ninth, our model was trained and evaluated exclusively at $T = 3$. Whether the self-correction mechanism converges to a stable attractor at longer horizons ($T \gg 3$) or eventually degrades through feedback accumulation (an "overthinking" effect) remains untested.

Tenth, our decoupled neuron importance analysis (Section 3.5) assumes that zeroing recurrent weights cleanly isolates the correction contribution. However, removing recurrent input shifts the pre-activation mean at Hidden Layer 1, potentially pushing some neurons into inactive regimes due to the ReLU nonlinearity, which may partially confound the separation of feedforward and recurrent contributions.

Eleventh, our primary statistical tests rely on N=10 independent models to maintain an exact nonparametric testing regime. While this avoids large-sample approximations, it bounds the minimum achievable p-value (approximately 0.002 before correction), limiting the statistical power to distinguish subtle effect size differences between conditions—for example, between C1 and C2. Additionally, standard non-parametric bootstrap confidence intervals on N=10 may exhibit undercoverage due to the sparsity of the empirical distribution; these CIs should be interpreted as secondary to the exact Wilcoxon p-values.

5 Conclusion

In a network of 35 neurons, self-correction develops when learning conditions permit temporal exploration and gradient flow through the output-feedback pathway. Once developed, self-correction depends entirely on the recurrent loop: ablating it eliminates correction while preserving baseline performance. Non-veridical feedback—whether statistical noise (C1) or the well-formed output of an identically trained clone (C2)—is more harmful than no feedback at all, with in-distribution clone feedback producing comparable degradation. C2 rules out marginal-distribution mismatch and demonstrates learned feedback-contract specificity: the recurrent weights are co-adapted to the model's own output geometry, and substituting another model's equally valid output reverses the sign of correction. Parameter-matched (D′) and compute-matched (D″) feedforward controls confirm that neither additional parameters nor computational depth can produce iterative correction without feedback content. These findings hold across 68% of 80 hyperparameter configurations within the explored $w_1/w_2/\tau$ grid.

Our results provide suggestive empirical support for the claim that, at least in minimal recurrent architectures, output-feedback processing and self-correction are not cleanly separable (Kirsch & Schmidhuber, 2022), offering ablation evidence from a fully transparent minimal model. The "fake mirror" effect—showing

that non-veridical feedback is worse than none, and that in-distribution impostors are at least as damaging as random noise—is, to our knowledge, novel, and opens a new diagnostic approach for studying output-feedback processing in neural systems.

We note, as a closing reflection, that these results point to a principle worth testing beyond toy models: the capacity for self-correction may not be cleanly separable from the capacity for output-feedback processing.

Data Availability

All code, data, and analysis scripts are publicly available at `https://github.com/softkorea/recurrent-self-correction-ablation`. The experiment is implemented in pure NumPy and can be reproduced on a standard desktop computer in approximately one hour. Raw experimental data (3,960 primary rows + 800 sweep rows) are included as CSV files.

Author Contributions

Sungmoon Ong conceived the hypothesis, designed the experiment, implemented the code, conducted the analyses, and wrote the manuscript. The experimental design was iteratively refined through structured review sessions with multiple AI systems, whose contributions are acknowledged below.

Acknowledgments

The experimental design benefited from structured review by AI systems that identified critical methodological improvements: the static-input task design that separates memory from self-reference, the scrambled feedback control group, the clone feedback control group that controls for marginal-distribution mismatch, multi-seed validation, temperature scaling for gradient flow, time-weighted loss for emergence induction, parameter-matched feedforward control, structured ablation control, noise-level sweep, and statistical design refinements including bootstrap confidence intervals and multiple comparison correction. The author takes full responsibility for all scientific claims and interpretations.

Conflict of Interest

The author declares no conflict of interest. This research was conducted independently without institutional or corporate funding.

References

Andrea Banino, Jan Balaguer, and Charles Blundell. Pondernet: Learning to ponder. *arXiv preprint arXiv:2107.05407*, 2021.

Stephen M Fleming and Raymond J Dolan. The neural basis of metacognitive ability. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 367(1594):1338–1349, 2012.

Alex Graves. Adaptive computation time for recurrent neural networks. *arXiv preprint arXiv:1603.08983*, 2016.

Louis Kirsch and Jürgen Schmidhuber. Eliminating meta optimization through self-referential meta learning. *arXiv preprint arXiv:2212.14392*, 2022.

Ji-An Li, Marcus K Benna, and Marcelo G Mattar. Discovering cognitive strategies with tiny recurrent neural networks. *Nature*, 644(8078):993–1001, 2025.

Haiping Ma, Wenlian Lu, and Jianfeng Feng. Spontaneous emergence of metacognition in neuronal computation. *Physical Review Research*, 7(3):033188, 2025.

Jorge Morales, Hakwan Lau, and Stephen M Fleming. Domain-general and domain-specific patterns of activity supporting metacognition in human prefrontal cortex. *Journal of Neuroscience*, 38(14):3534–3546, 2018.

Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets. In *International Conference on Learning Representations*, 2022.

Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, et al. Emergent abilities of large language models. *Transactions on Machine Learning Research*, 2022a.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, volume 35, 2022b.

Yusuke Yamato, Reiji Suzuki, and Takaya Arita. Evolution of metamemory based on self-reference to own memory in artificial neural network with neuromodulation. *Scientific Reports*, 12(1):6233, 2022.

A Additional Figures

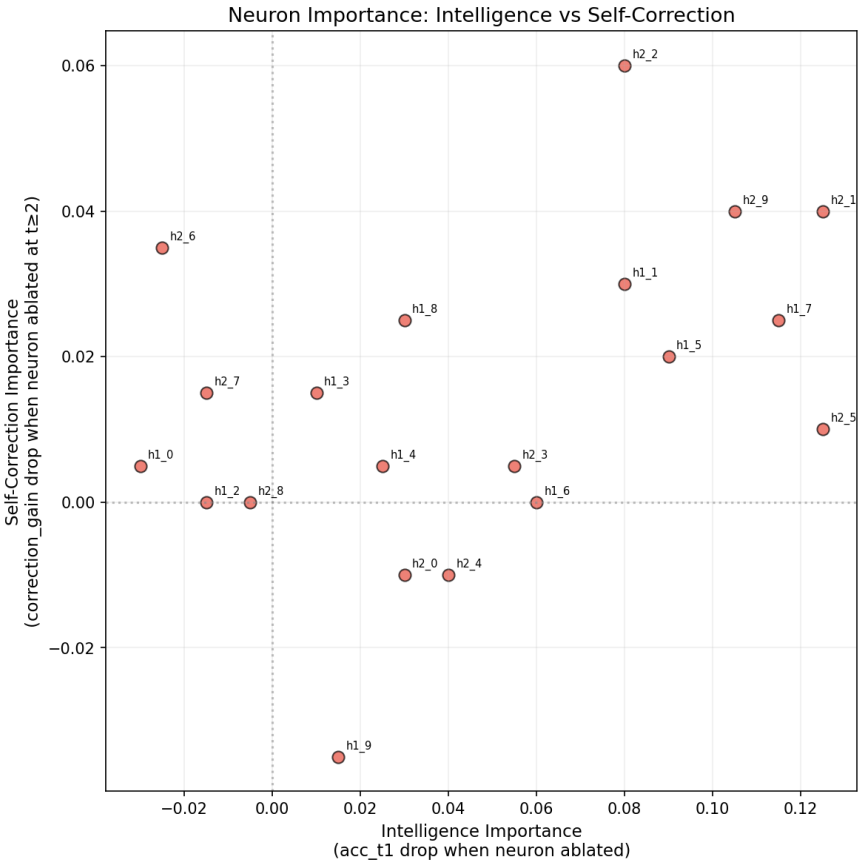

Figure 4: Neuron importance: intelligence (feedforward) vs. correction (recurrent) contribution. Exploratory analysis on a single model (seed=0).

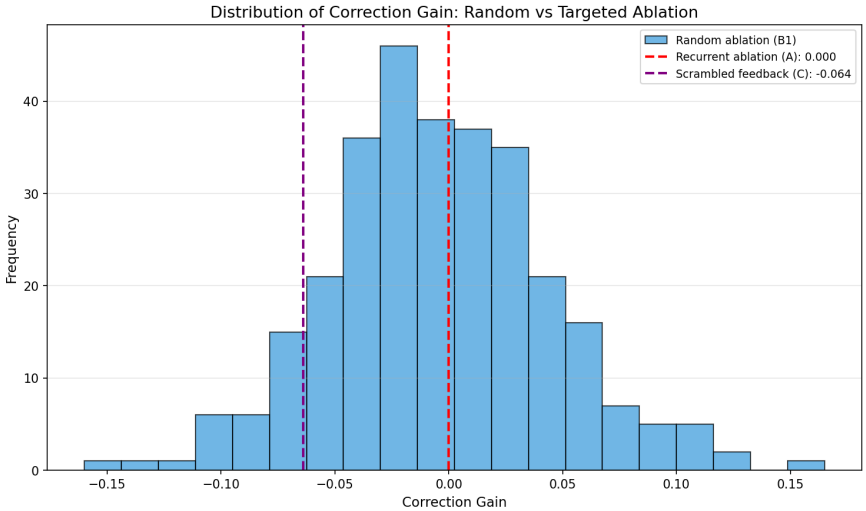

Figure 5: Distribution of correction gain under random ablation (Group B1, 30 repeats $\times$ 10 models). Vertical lines show Group A and Group C1 positions for comparison.